

Fourth Year Project Logbook

Contents

| | |
|---|-----------|
| Week 1 | 4 |
| 30/09/13 - Exploring simple case with PAM modulation | 4 |
| 01/10/13 - Implementing Raised Cosine functions | 5 |
| 02/10/13 - Wrapping Up the Initial PAM Model | 6 |
| 03/10/13 - Delving deeper into Mathematica | 6 |
| Week 1 Summary | 6 |
| Goals for Week 2 | 7 |
| Week 2 | 7 |
| 07/10/13 - Matrix manipulations | 7 |
| 08/10/13 - Fixed I.R. and Kernel Density Estimation | 7 |
| 09/10/13 - Setting up Digital Comms Lab PC | 9 |
| 10/10/13 - Probing the Elec Eng network | 9 |
| Week 2 Summary | 11 |
| Goals for Week 3 | 11 |
| Week 3 | 11 |
| 14/10/13 - Running longer scripts on the EDA machines | 11 |
| 15/10/13 - Reducing output size | 11 |
| 16/10/13 - Moving onto 4-PAM | 12 |
| Week 3 Summary | 12 |
| Goals for Week 4 | 13 |

| | |
|--|-----------|
| Week 4 | 13 |
| 21/10/13 - Implementing the Gram-Charlier series | 13 |
| 22/10/13 - More Gram-Charlier series | 14 |
| 23/10/13 - Proper Gram-Charlier plots | 14 |
| Week 4 Summary | 14 |
| Goals for Week 5 | 14 |
| Week 5 | 14 |
| 29/10/13 - Comparing Gram-Charlier to Simulation | 14 |
| 30/10/13 - Applying the Tikhonov Distribution | 15 |
| Week 5 Summary | 18 |
| Goals for Week 6 | 18 |
| Week 6 | 18 |
| 04/11/13 - Fixing errors | 18 |
| 05/11/13 - Corrected simulation results | 18 |
| Week 6 Summary | 18 |
| Goals for Week 7 | 19 |
| Week 7 | 19 |
| 11/11/13 - Returning to the UNIX machines | 19 |
| 19/11/13 - Day 9 of Week 7 | 19 |
| Week 7 Summary | 19 |
| Goals for Week 8 | 19 |
| Week 8 | 21 |
| Week 8 Summary | 21 |
| Goals for Week 9 | 21 |
| Week 9 | 21 |
| 13/01/14 - Examining combining and fading | 21 |
| 14/01/14 - Planning an analytical analysis of the non-fading case . . . | 22 |
| 15/01/14 - Evaluating the Gram-Charlier distribution over the Tikhonov distribution | 22 |

| | |
|---|-----------|
| 16/01/14 - Results for the above | 22 |
| Week 9 Summary | 22 |
| Goals for Week 10 | 24 |
| Week 10 | 24 |
| 20/01/14 - Rayleigh Fading Results | 24 |
| 21/01/14 - Implementing Dave's suggestions | 24 |
| Week 10 Summary | 25 |
| Goals for Week 11 | 25 |
| Week 11 | 25 |
| 27/01/14 - Simulating EGC and MRC | 25 |
| 30/01/14 - EGC vs MRC results | 25 |
| 01/02/14 - SNR increase to 28dB | 25 |
| Week 11 Summary | 27 |
| Goals for Week 12 | 27 |
| Week 12 | 27 |
| 03/02/14 - Presentation work | 27 |
| 06/02/14 - Presentation test run and rework | 27 |
| Week 12 Summary | 27 |
| Week 13 | 29 |
| 11/02/14 - Analytical description of non-fading case | 29 |
| 12/02/14 - Viability of independent variate route | 29 |
| 13/04/14 - MRC adaptation of Gram-Charlier distribution | 30 |
| 14/04/14 - Explaining the oscillations | 30 |
| Week 13 Summary | 30 |
| Goals for Week 14 | 30 |

| | |
|---|-----------|
| Week 14 | 31 |
| 16/02/14 - Bit Error Rate Calculation | 31 |
| 18/02/14 - Averaging BER over Tikhonov and Rayleigh distributions . | 31 |
| 19/02/14 - Root-finding tweaks | 31 |
| 20/02/14 - Discussion of Gram-Charlier accuracy | 32 |
| 21/02/14 - Overcoming the timing error problems | 32 |
| 22/02/14 - Further problems for high timing offsets | 32 |
| Week 14 Summary | 32 |
| Goals for Week 15 | 33 |
| Week 15 | 33 |
| 24/02/14 - Root-finding-defying wiggles | 33 |
| 25/02/14 - Correct optimum DRB measurements | 33 |
| 26/02/14 - Integration | 34 |
| Week 15 Summary | 35 |
| Goals for Week 16 | 35 |
| Week 16 | 35 |
| 03/03/14 - Rewriting analytical implementation | 35 |
| 04/03/14 - Strengthening the implementation | 36 |
| 05/03/14 - Assessing the accuracy of the implementation | 36 |
| Week 17 | 36 |
| 14/03/14 - Pre-open day progress | 36 |
| Week 18 | 37 |
| 17/03/14 - Comparing Analytic and Simulation results side-by-side . . | 37 |
| 18/03/14 - Fixing the Gram-Charlier code | 37 |

Week 1

30/09/13 - Exploring simple case with PAM modulation

I received the *PAM.pdf* file outlining the case where a signal is sent through a channel with AWGN and received with a timing error at the receiver. I

read through the file several times to get an understanding of the underlying equations.

Leaving the Gram-Charlier series aside for the moment, I started getting to grips with Mathematica and implementing the transmission system model:

$$X = \omega_0 g_0 + \sum_{k=1}^{40} (\omega_{-k} g_k + \omega_k g_k) + \nu$$

where $g_k = g((\Delta + k)T)$, $g(t) = (u_T * h_l * u_R)(t) \times \cos(\theta)$ and ν is a zero-mean Gaussian random variate with $\sigma_\nu^2 = N_0 \varepsilon_R$.

I learnt the basics of the interface, and began implementing the filter and channel impulse responses (I.R.). I need to double-check the definition of the Root-Raised Cosine (RRC) Filter, as the impulse response wasn't as expected.

Later, I found the [correct form for the RRC](#) and double-checked it using Octave. The equation used is listed below. A plot showed that this equation is invalid at $t = \left[-\frac{T_s}{4\beta}, 0, \frac{T_s}{4\beta}\right]$, so I plan to find its limit at these points using Mathematica to obtain the complete solution.

$$h_{RRC}(t) = \frac{2\beta}{\pi\sqrt{T_s}} \frac{\cos\left[(1+\beta)\frac{\pi t}{T_s}\right] + \frac{\sin\left[(1-\beta)\frac{\pi t}{T_s}\right]}{\frac{4\beta t}{T_s}}}{1 - \left(\frac{4\beta t}{T_s}\right)^2}$$

01/10/13 - Implementing Raised Cosine functions

I implemented the function above in Mathematica, and using the `Limit` function found the value of the function at the following undetermined points:

$$h_{RRC}(t) = \begin{cases} \frac{4\beta + \pi(1-\beta)}{2\pi\sqrt{T_s}} & t = 0 \\ \frac{\beta}{2\pi\sqrt{T_s}} \left(\pi \sin\left[\frac{(1+\beta)\pi}{4\beta}\right] - 2\cos\left[\frac{(1+\beta)\pi}{4\beta}\right] \right) & t = \pm \frac{T_s}{4\beta} \\ \frac{2\beta}{\pi\sqrt{T_s}} \frac{\cos\left[(1+\beta)\frac{\pi t}{T_s}\right] + \frac{\sin\left[(1-\beta)\frac{\pi t}{T_s}\right]}{\frac{4\beta t}{T_s}}}{1 - \left(\frac{4\beta t}{T_s}\right)^2} & \text{otherwise} \end{cases}$$

I also implemented the Raised Cosine function for the channel function, using the impulse response below¹. I was unable however to convolve the receiver and

¹Proakis, "Digital Communications"

transmitter filter functions using the `Convolve` function, even when I limited the impulse response using a `UnitBox`.

$$h_{RC}(t) = \frac{\text{sinc}\left(\frac{\pi t}{T}\right) \cos\left(\beta \frac{\pi t}{T}\right)}{1 - \left(2\beta \frac{t}{T}\right)^2}$$

I looked into Mathematica's treatment of the Gaussian distribution, and figured out how to generate random noise vectors following a Gaussian distribution, as well as how to generate a list of random binary symbols.

After discussing the convolution issue with David, he suggested that the channel should be initially modelled as ideal and therefore the overall channel and filter I.R. $g(t)$ can be defined as a Raised Cosine function, as defined above. I should therefore be ready to implement the simple ISI model tomorrow.

02/10/13 - Wrapping Up the Initial PAM Model

I pulled together the Raised Cosine function and random number generator to implement the given simplified function for the PAM receiver output, given below. Playing around with the settings, I was able to show how the g_k function increases with the timing error. I decided to study the Mathematica environment a little more before carrying on with any programming.

$$X = \omega_0 g_0 + \sum_{k=1}^{40} (\omega_{-k} g_{-k} + \omega_k g_k) + \nu$$

03/10/13 - Delving deeper into Mathematica

I devoted some time into looking through Michael Quinlan's notebooks and better understanding the workings of the `Table` functions and the various plotting options. Fortunately my notebook was corrupted so I was able to rewrite it and understand the model a bit more. I need to figure out what variance value the noise PDF should take on, as the noise appears to be overwhelming the timing error effects. Translating the resulting PDF's into patterns is another question that needs some thought.

Week 1 Summary

Week 1 was mostly spent becoming acquainted with Mathematica and getting a feel for the equations underlying PAM transmissions. A simple model of a PAM receiver was constructed.

Goals for Week 2

- The PAM model will need to be extended to calculate the optimum decision region boundary from the estimated PDF.
- A better setup will be required to perform large-scale simulations within an acceptable time period. We will look into applying for an account on the Boole cluster.

Week 2

07/10/13 - Matrix manipulations

I decided to spend another day learning about the Mathematica environment, in particular matrix manipulation and generation. I looked into the **Apply**, **Map** and **Partition** functions and wrote some examples to figure out how to convert mathematical problems to Mathematica notation using matrices. I hope to convert the code to use matrices tomorrow to hopefully simplify and speed things up.

I also implemented David's equation for properly calculating the AWGN function variance from SNR², from last Friday's meeting.

08/10/13 - Fixed I.R. and Kernel Density Estimation

The first job was to rewrite the code to make use of the simple dot operator to calculate all the ISI components³. With the new code I was able to carry out many more runs and get much more detailed output. In addition, when I was rewriting the code I noticed a typo in the Raised Cosine I.R. that was degrading performance in the perfectly synchronised case. With both of these changes made, I decided to use Kernel Density Estimation to see what effects the timing offset has.

Using offsets of 10^{-15} , 0.05, 0.1 & 0.15, the following values of $g_k, k \in \{-40 \dots -1, 1 \dots 40\}$ were calculated.

²See *davenotes.pdf*

³The ISI components are now calculated using:

$$\left[\sum_{k=0}^{k=40} (g_k \omega_k^1 + g_{-k} \omega_{-k}^1) \cdots \sum_{k=0}^{k=40} (g_k \omega_k^m + g_{-k} \omega_{-k}^m) \right] = [g_{-40} \cdots g_{-1} g_1 \cdots g_{40}] \bullet \begin{bmatrix} \omega_{-40}^1 & \cdots & \omega_{-40}^m \\ \vdots & & \vdots \\ \omega_{-1}^1 & \cdots & \omega_{-1}^m \\ \omega_1^1 & \cdots & \omega_1^m \\ \vdots & & \vdots \\ \omega_{40}^1 & \cdots & \omega_{40}^m \end{bmatrix}$$

where ω_k^j is the k 'th ISI with the j 'th timing offset.




Figure 1: g_k linear plot




Figure 2: g_k log plot

Using `SmoothKernelDistribution` to perform Kernel Density Estimation with 1 million points produced the following estimated PDFs for both possible transmitted values. As the timing error increases, we note that the PDF spreads out, but the mean remains steady.




Figure 3: kernel density estimation $\omega_0 = 1$

09/10/13 - Setting up Digital Comms Lab PC

With Ger's help, I set up an account on **Digital Comms Lab 1** & **Digital Comms Lab 2** and got the internet working. Mathematica 8 is installed and working on both machines, we will have to consider whether an upgrade to Mathematica 9 would be useful or not. Git and VNC or similar have to be installed next. A request was made to the Boole cluster for access for this machine, however the email given (bcrisupport@bcricri.ucc.ie) was invalid.

10/10/13 - Probing the Elec Eng network

After finding out the Boole cluster was no more, I used today to examine what hardware I had available to me. I got access from Ger to the public UEPC004 server, and from there I am able to access machines on the elec eng network. I set up a *Remote Desktop Protocol* link to **Digital Comms Lab 1** through this server, allowing me to control the machine from any location. I am also able to log remotely into EDA lab machines, and run Mathematica 6 on those machines.⁴

⁴The GUI does not work when using `ssh` to access the EDA Lab machines, but using the command `math` to start and operate Mathematica kernels does.




Figure 4: kernel density estimation $\omega_0 = -1$

Ger has been known to tweak machines in response to personal requests, so if asked nicely he may let me use two or three of these machines concurrently.

Given these resources, I feel there are three ways I could continue:

- I could upgrade to the latest version of Mathematica on all machines, and set up a Mathematica cluster with **Digital Comms Lab 1** as the front end and the EDA Lab PCs as remote nodes. With this setup, all machines would act as one (as in a traditional cluster). This would be the easiest to use, but would require considerable work to set up.
- I could use the **MathLink** interface to achieve a similar, lower-level version of the former, with the EDA Lab machines as independent, remote slaves and **Digital Comms Lab 1** sending commands to these slaves and collating the replies. This setup is distributed computing with a star topology, and would be easier to setup. The downside is that the code needs to manually divide the task between each of the nodes, and needs to be well designed to minimise network delays.
- I could simply run the code in parallel on each of the machines available to me, dumping the results to text files, and collate the data at the end. This would require no setup, and code written on any machine would only require porting to another version of Mathematica. Additionally this seems like it would deal best with hiccups such as machines going down and it does not depend on a connection between the machines. The downside is there would be some overhead with collecting the results afterwards.

Week 2 Summary

I fixed the code written last week and began setting up my simulation environment.

Goals for Week 3

- Work out a setup that will allow me to carry out large-scale simulations.
- Adapt the previous code to run in parallel and produce useful machine-readable output.

Week 3

14/10/13 - Running longer scripts on the EDA machines

Today I spent some time figuring out how to build and run scripts on the EDA machines. I found that defining a module in a text file and copying the Mathematica code into the module allows the code to be called with input arguments, and writing the output to a text file and placing the module in a loop allows each pass to be recorded for later parsing⁵. After running the code overnight, this system appears to work, and is scaleable over multiple machines. The main disadvantage is the size of these files (7.7MB per 400,000 values), so I must either figure out how to transfer them over the network or see whether reducing the precision of the output values will reduce the file sizes.

15/10/13 - Reducing output size

Given the 15GB of samples produced the night before was far too much to pull off the machine, I copied 20 million of the samples and plotted them to make sure the script had worked in practice⁶. I then looked into how I could reduce the size of the output produced, and decided to replace the `SmoothKernelDistribution` function (which came in after Mathematica 6.0 and therefore couldn't be used on the EDA machines) with a fine-grained histogram function⁷. This allowed me to add the probabilities generated in each sweep to those generated before and keep the output to a handful of 1kB files. I ran the simulation overnight to check it.

⁵Using the `Get` and `Put` methods. The `DumpSave` method is supposed to be more efficient, but was added after Mathematica 6.

⁶For the record, I could only use a fraction of them, as loading all 20 million samples crashed the machine for over an hour.

⁷I am assuming that both approach the true PDF as $N \rightarrow \infty$

16/10/13 - Moving onto 4-PAM

Checking the output from the night before, I get a similar PDF plot as with the `SmoothKernelDistribution` function. I therefore modified the code to examine all 3 decision region boundaries in a 4-PAM system and ran the simulation for 100 million samples per condition. The resulting distributions shown below show increased probability of error with timing error, as expected, but decision region boundaries in this case remain the same.




Figure 5: PDF for 4-PAM, $\omega_0 \in -3, -1, 1, 2, 10^8$ samples

I could imagine finding a value for the probability of error and moving onto PSK systems as the next steps in the process.

Week 3 Summary

Code was written that could be executed in parallel on multiple machines, and this was demonstrated in practice. The code was extended to the 4-PAM case, and showed no change in optimum decision region boundaries. Upon later consultation with Dave, it seems this is because the decision region boundaries shift due to a change in the g_0 term, and not the appearance of ISI components due to the g_k terms; the latter was believed to be the expected cause, and so the g_0 term was assumed to be 1 in the code.

Goals for Week 4

- Re-run the simulations to see if implementing the change in g_0 with timing error changes the location of the optimum decision region boundaries.
- If so, it would be interesting to see if the Gram-Charlier approximation produces the same boundary locations.

Week 4

21/10/13 - Implementing the Gram-Charlier series

Over the weekend, I implemented the g_0 term fix discussed in our Friday weekly meeting and re-ran the simulation, this time across two machines. Results showed that the Decision Region Boundary is displaced towards the origin as the timing offset increases.




Figure 6: PDF for 4-PAM, $\omega_0 \in \{-3, -1, 1, 2\}$

I spent Monday carrying out two tasks:

1. I re-wrote Dave's Gram-Charlier equations⁸ for Mathematica, and should

⁸See *davenotes.pdf*

be ready to try them out tomorrow.

2. I modified the PAM simulation with a coarser-grained histogram, but more timing offset values, in order to see how the decision variate varies with timing offset. The results should be available in the morning.

22/10/13 - More Gram-Charlier series

The simulation results showed that the decision region boundaries did decrease with timing error, however the histogram was not fine-grained enough to accurately determine the exact boundary locations, so the simulation was re-run with more bins.

I fixed some bugs in my implementation of the Gram-Charlier series and was able to generate a few plots, which were very similar to those generated by the simulator, albeit with half the amplitude. A goal for tomorrow is to generate the plots with identical timing offsets to the simulation and compare both plots.

23/10/13 - Proper Gram-Charlier plots

The simulation results had been appended to the previous set of results by accident, so the whole thing had to be run again for tomorrow. On a more positive note, I noticed a missing power in my implementation of the Gram-Charlier series, and the plots are now a lot closer to those generated previously.

Week 4 Summary

I implemented the Gram-Charlier series and was able to compare the results from the simulation and the Gram-Charlier series. These are close, but not exact, so we will have to look closely at where the differences may be coming from.

Goals for Week 5

- I will make use of the long weekend to run some extra-long simulations and compare these to the Gram-Charlier series.

Week 5

29/10/13 - Comparing Gram-Charlier to Simulation

The simulations ended, and I was able to compare simulated and gram-charlier PDF plots. I extracted a rough estimate of the decision region boundaries given

by both methods and compared them to the corresponding values of $2g(\Delta)$, and found very close correlation.

`{1.99962, 1.99062, 1.96267, 1.91668, 1.85355}`




Figure 7: Gram Charlier approximation of $P(\omega_0 = 1, R) - P(\omega_0 = 3, R)$

30/10/13 - Applying the Tikhonov Distribution

I was able to implement the Tikhonov Distribution using the equation provided in *PAMTikhonov.pdf*:

$$F_{\Delta}(y) = \frac{\text{Exp} \left[\frac{\cos(2\pi y)}{(2\pi\sigma_{\Delta})^2} \right]}{I_0 \left(\frac{1}{(2\pi\sigma_{\Delta})^2} \right)} \text{ where } -\frac{1}{2} \leq y \leq \frac{1}{2}$$

Given these timing error probabilities and the optimum decision region boundaries for each timing error, I calculated the overall optimum decision region boundary for each timing error probability distribution using

$$B_{\text{OPT}} \sim \sum_{\Delta} P(\Delta) B_{\text{OPT}, \Delta}$$

It is important to note that with increasing variance, the probability density function places more weight on larger timing errors outside the range simulated, so these results are less accurate for higher variances.




Figure 8: Simulation of $P(\omega_0 = 1, R) - P(\omega_0 = 3, R)$, $N=3 \times 10^6$




Figure 9: Comparison of Gram-Charlier Decision Region Boundaries and $2g(\Delta)$ estimation ($0.01 \geq \Delta \geq 0.2$)




Figure 10: Optimum Decision Region Boundary for various timing error probability distributions

Week 5 Summary

In week 5, I calculated the optimum decision region boundary for a range of timing offsets, through simulation and the Gram-Charlier approximation. I demonstrated a close correlation between these boundaries and the $2g_k$ term. A slight difference between the Gram-Charlier approximation was found to be due to a typo in its implementation. I applied the Tikhonov distribution to the calculated optimum decision region boundaries for each timing offset, in order to calculate an optimum decision region boundary for a given Tikhonov distribution of timing offsets

Goals for Week 6

- On the simulation side, a key goal for week 6 is to randomly generate timing offsets according to the Tikhonov distribution and apply these to the simulation as timing offsets, in order to verify correlation with the Gram-Charlier and $2g_k$ approximations.
- A typo in the Gram-Charlier implementation has been found and corrected, and it would be interesting to see if this approximation matches $2g_k$.

Week 6

04/11/13 - Fixing errors

Dave took a look at my code and spotted errors which I fixed. The fixed Gram-Charlier implementation was found to match $2g_k$ very closely. The fixed simulation was left to run overnight; unfortunately Mathematica 6.0 running on the Unix machines was unable to run it, so the number of points had to be reduced.

05/11/13 - Corrected simulation results

The produced PDFs were too inaccurate to properly calculate the zero crossing points, so the simulation will have to be run over several days.

Week 6 Summary

A simulation was constructed that generated timing error offsets according to a Tikhonov distribution of predetermined variance, and used to produce received symbol PDFs. The simulation was found to run very slowly, and could only be run on Mathematica 9. Ger has been asked whether it would be possible to upgrade the Unix machines to this version and he will look into it.

Goals for Week 7

- Continue running the simulation, trying to speed it up if at all possible.

Week 7

11/11/13 - Returning to the UNIX machines

The UNIX machines were upgraded to Mathematica 9 over the weekend, so I was able to port the code in order to run off these. In addition, Dave suggested that I look into parralising the code. Since these were dual-core machines I was able to make use of Mathematica's `ParallelTable` function to reduce run times a little. Th simulation will have to run over several days, however, as the expected deviation in optimum decision region boundary is very small.

19/11/13 - Day 9 of Week 7

After several days of running the simulations, we found the optimum decision region boundaries given by the simulations, in red, converged to roughly those predicted by averaging the optimum decision region boundary of a timing offset over the tikhonov distribution of timing offsets, in blue, given by the equation:

$$B_{\text{OPT}} \sim \sum_{\Delta} P(\Delta) B_{\text{OPT}, \Delta}$$

Week 7 Summary

Simulations supported the theory that the optimum decision region boundary in the presence of statistically distributed receiver timing errors will decrease from the expected value. Additionally, it was shown through simulation that the new optimum decision region boundary can be aproximated, assuming a known distribution of these timing errors, by averaging the optimum decision region boundary given each timing offset over the distribution of timing offsets.

Goals for Week 8

- Verify that the Gram-Charlier series provides an adequate approximation to the received symbol PDF in the presence of timing errors.
- Provide numerical values for the probability of error P_e in the presence of a distribution of timing errors.




Figure 11: Optimum Decision Region Boundary for various timing error probability distributions

Week 8

Week 8 Summary

After spending the Christmas break reviewing the literature on Rayleigh fading, mainly *Proakis*, I implemented the Rayleigh distribution as a modified Nakagami-n distribution⁹ I implemented this as well as a multi-receiver combining system and added it to the existing simulation. Following a meeting with Dave and Colin, we decided to proceed as follows:

- Run simulations similar to those run before Christmas to determine numerically the optimum decision region boundaries and corresponding error probabilities in the presence of Rayleigh fading with Equal Gain Combining (EGC).
- Alongside this, attempt to provide an analytical derivation for the optimum decision region boundaries in the presence of timing errors in a non-fading environment. Should the above simulations show merit in modifying the decision region boundaries, this could then be adapted to the fading case later on.

Goals for Week 9

- Implement the simulation. This is expected to run for the duration of Week 8, and demonstrate the optimum decision region boundaries for an EGC receiver with Rayleigh fading. The main goal is to determine if different decision region boundaries would reduce the probability of error.
- Attempt to describe the system analytically, ignoring fading. This will be based on a Gram-Charlier approximation.

Week 9

13/01/14 - Examining combining and fading

Over the weekend, three separate approximate simulations were run to examine receiver performance in three cases:

1. Similar to before Christmas, a single-receiver system in a non-fading environment
2. A multi-antenna system using EGC in a non-fading environment
3. An EGC system in the presence of Rayleigh fading

⁹The Rayleigh distribution is a special case of the Nakagami-n distribution, with $n=1$

Initial simulations for case 3. showed reduced optimum decision region boundaries in the presence of Rayleigh fading, which appeared to remain constant with changing timing offset variance. A simulation of case 2. showed changing optimum decision region boundaries for different timing offset variances. A more detailed simulation of case 3. examining more variances was commenced, and is expected to finish mid-week.

14/01/14 - Planning an analytical analysis of the non-fading case

I found an implementation of the Gram-Charlier series and Tikhonov distribution I had developed previously that could be used as a basis for an analytical exploration of the project topic. The code produces a Gram-Charlier PDF for a range of timing offsets, calculates the optimum decision region boundary for each, and averages these boundaries over the Tikhonov distribution. Dave suggested that while this approach wasn't mathematically correct, if the Gram-Charlier PDF's were averaged over the Tikhonov distribution to provide an overall PDF, the optimum decision region boundary could be calculated from this.

15/01/14 - Evaluating the Gram-Charlier distribution over the Tikhonov distribution

I carried out the changes detailed above, and added a loop to numerically estimate the location of the PDF crossings, thereby estimating the optimum decision region boundaries. Unfortunately there was not enough time to run the code over all possible conditions.

One of the simulations begun at the start of the week quit unexpectedly, and had to be restarted, moving back the expected end-date for the simulations to Friday.

16/01/14 - Results for the above

I was able to run the Tikhonov-Gram-Charlier code described above and thus plot the optimum decision region boundaries for Tikhonov-distributed timing offsets, assuming a Gram-Charlier approximation for the received symbol. These correlated strongly with the optimum boundaries found through simulation.

Week 9 Summary

A simulation covering the 4-PAM case studied before Christmas was extended with a 4-channel EGC system and simulated Rayleigh channel fading. This simulation was run over the course of the week and (*proved xyz...*)




Figure 12: Plot of optimum decision region boundaries versus Tikhonov-distributed timing error variance, from Gram-Charlier approximation and simulation respectively

A rough analytical study of the non-fading timing error model examined before Christmas was carried out. The PDF's of various cases with a fixed timing offset were generated, and averaged over a discretised Tikhonov distribution to estimate the PDF of the system in the presence of Tikhonov distributed timing offsets. This was carried out for a range of Tikhonov distribution variances to examine different conditions. The optimum decision region boundary for each condition was found numerically and found to correlate strongly with the locations determined through simulation previously, with the exception of some low-variance points.

Goals for Week 10

- Examine the performance of the system in the presence of Rayleigh distributed channel fading.
- Determine the probability of error of the system in the non-fading case, and evaluate any performance gains from using the calculated optimum decision region boundaries.

Week 10

20/01/14 - Rayleigh Fading Results

Simulation ended over the weekend and showed an optimum decision region boundary around 1.72, for Tikhonov variances of 0.001 to 0.010. This is an interesting result, as it suggests that in this particular case the optimum decision region boundary is only loosely related to the Tikhonov variance, which would imply that knowledge of the timing statistics is less important in implementation than previously thought.

Dave produced an analytical solution for the optimum decision region boundary in the non-fading case, which demonstrated clear BER gains (20-33%).

21/01/14 - Implementing Dave's suggestions

Following some consideration, Dave suggested that the SNR of 8dB used in the simulation was too low for 4-PAM, and it's possible that the ODRB of 1.72 seen was the lower bound for the ODRB. We settled on a more reasonable SNR of 20dB, and decided to apply each set of random conditions to both possible sent symbols $\omega_0 = 1,3$ so that we would only have to run the simulations once. Another suggestion made by Dave was to apply a different timing offset to each channel, as each channel has a separate receiver, and therefore independent timing. I also implemented Maximal-Ratio Combining as an option in order to compare the performance of both systems.

Week 10 Summary

This week we were able to show reduced optimum decision region boundaries in the presence of Rayleigh fading with Equal-Gain Combining. These results showed that the performance of the EGC receiver could be considerably improved when facing low SNR. It was decided that the short-term goal would be to evaluate the receiver's performance with a more reasonable SNR, and compare it to the performance of a Maximal-Ratio Combining system under similar conditions.

Work stopped mid-week to facilitate the author's funding application, and is expected to resume in full next week.

Goals for Week 11

- Evaluate the error rate of an EGC receiver with optimised decision region boundaries in Rayleigh fading with an SNR of 20dB.
- Similarly, evaluate the error rate of a MRC receiver in similar conditions.
- Compare and contrast the performance of each, taking into account the higher area and power requirements of the latter.

Week 11

27/01/14 - Simulating EGC and MRC

The changes discussed in the last entry were implemented, and both simulations started.

30/01/14 - EGC vs MRC results

These simulations took a little longer than expected due to the extra processing required for MRC. Error rates were found to decrease after redefining the decision region boundary of the EGC receiver. The MRC system suffered more from the effects of the timing error, but the reduced EGC error rates were still a good bit higher than the equivalent error rates using MRC. Nonetheless, a 20-33% improvement was found over the un-tweaked decision region boundaries.

Redefining the decision region boundary for the MRC also showed improvements, although these were a more modest 7-20%.

01/02/14 - SNR increase to 28dB

In order to determine how much of an effect noise has on the system in the presence of fading, the simulation was restarted with an SNR of 28dB.

| | | | | | | | | | | |
|----------|-----------------------|-----------|----------|----------|---------|---------|---------|--------|--------|--------|
| Tikh var | 0.001 | 0.002 | 0.003 | 0.004 | 0.005 | 0.006 | 0.007 | 0.008 | 0.009 | 0.01 |
| Prev ER | 0.111 | 0.115 | 0.119 | 0.124 | 0.129 | 0.133 | 0.139 | 0.148 | 0.154 | 0.159 |
| New ER | 0.0751 | 0.0801 | 0.0844 | 0.0900 | 0.0951 | 0.0994 | 0.106 | 0.115 | 0.122 | 0.126 |
| % drop | 33. | 30. | 29. | 27. | 26. | 25. | 24. | 22. | 21. | 20. |
| MRC ER | 6.00×10^{-8} | 0.0000298 | 0.000354 | 0.00121 | 0.00320 | 0.00565 | 0.00943 | 0.0146 | 0.0193 | 0.0265 |
| MRC ER | 8.00×10^{-8} | 0.0000159 | 0.000218 | 0.000970 | 0.00246 | 0.00470 | 0.00818 | 0.0129 | 0.0179 | 0.0242 |
| % drop | -33. | 47. | 38. | 20. | 23. | 17. | 13. | 12. | 7. | 9. |




Figure 13: EGC and MRC error rates, SNR=20dB

Week 11 Summary

The simulations ran this week proved that an error rate reduction could be obtained by merely redefining the location of the decision region boundaries, in the Rayleigh fading case. Although the EGC system could not be made to perform with comparable performance to the MRC system, we demonstrated rapidly deteriorating MRC performance with more severe timing offsets, vindicating our study of the issue.

Goals for Week 12

- Examine other conditions to gain a more general understanding of the effects of timing offsets in Rayleigh fading channels.
- Start on the analytical examination of timing offset.
- Compile a presentation for the seminars.

Week 12

03/02/14 - Presentation work

The results from the 28dB run showed very similar error rates, confirming my suspicion that the effects of fading are dominating over the effects due to AWGN seen in the non-fading case. A simulation at 12dB was started to see if this remained true at lower SNR's.

The seminar presentation was also compiled.

06/02/14 - Presentation test run and rework

The presentation was given a test-run, leading to improvements in the clarity of the descriptions.

The simulation ended, but unfortunately it turned out that the EGC simulation had been ran with an SNR of 8dB, so had to be run again.

Dave announced he has been able to start the analytical proof of the timing offset, and we should be able to hear from him at the start of next week.

Week 12 Summary

A relatively quiet week, with a focus on the seminar presentation.

| | | | | | | | | | | |
|----------|-----------------------|-----------|----------|----------|---------|---------|---------|--------|--------|--------|
| Tikh var | 0.001 | 0.002 | 0.003 | 0.004 | 0.005 | 0.006 | 0.007 | 0.008 | 0.009 | 0.01 |
| Prev ER | 0.111 | 0.115 | 0.119 | 0.123 | 0.129 | 0.134 | 0.140 | 0.148 | 0.153 | 0.159 |
| New ER | 0.0751 | 0.0797 | 0.0847 | 0.0883 | 0.0954 | 0.100 | 0.108 | 0.116 | 0.119 | 0.126 |
| % drop | 33. | 31. | 29. | 28. | 26. | 25. | 23. | 21. | 22. | 20. |
| MRC ER | 4.00×10^{-8} | 0.0000185 | 0.000286 | 0.000937 | 0.00287 | 0.00534 | 0.0101 | 0.0131 | 0.0188 | 0.0282 |
| MRC ER | 8.00×10^{-8} | 0.0000109 | 0.000168 | 0.000776 | 0.00219 | 0.00446 | 0.00886 | 0.0121 | 0.0172 | 0.0245 |
| % drop | -100. | 41. | 41. | 17. | 23. | 16. | 12. | 8. | 8. | 13. |




Figure 14: EGC and MRC error rates, SNR=28dB

Week 13

11/02/14 - Analytical description of non-fading case

Dave provided a proof for the Gram-Charlier approximation of the non-fading case, along with some accompanying code. Used the Gauss-Legendre method to combine the Gram-Charlier solution for a fixed timing offset with the Tikhonov timing offset distribution to provide a solution for the PDF at the input to the receiver, which is too long to reproduce here in full, but is approximately given by:

$$\frac{b-a}{2} \sum_{i=1}^N \omega_i T\left(\frac{b-a}{2} z_i - \frac{b+1}{2}\right) GC\left(\frac{b-a}{2} z_i - \frac{b+1}{2}\right)$$

where $\omega_i = \frac{2}{(1-z_i)(P(z_i))^2}$ and z_i are the roots of the Legendre polynomial $P_n(z)$.

I spent some time examining the Mathematica code he used to examine the above, and decided as an introduction to try to extend it to the Equal-Gain coming case. I used the property that for the PDF of the sum of two independent variants is the convolution of each variant's PDF (ie. for two variants X_1 and X_2 with PDF $f_1(x)$ and $f_2(x)$, the PDF of $Y = X_1 + X_2$ is $f_Y(x) = f_1(x) * f_2(x)$) to express the PDF of an m+1 antenna EGC system as

$$f_{EGC}(x) = f^{m*}(x)$$

where $f(x)$ is the PDF at each antenna and $f^{2*}(x)$, for example, is the double convolution $f(x) * f(x) * f(x)$.

I was then able to use the Gauss-Legendre method to describe this as an m-fold sum:

$$f_{EGC}(y) = f^{m*}(y) \simeq \sum_{a_1=1}^{n_f} \sum_{a_2=1}^{n_f} \cdots \sum_{a_m=1}^{n_f} 3^m \omega_{a_1} \omega_{a_2} \cdots \omega_{a_m} \\ f(3z_{a_m} - 3z_{a_{m-1}}) f(3z_{a_{m-1}} - 3z_{a_{m-2}}) \cdots f(3z_{a_2} - 3z_{a_1}) f(3z_{a_1} + 2) f(y - 3z_{a_m} - 2)$$

I was unfortunately unable to test this code before end of day.

12/02/14 - Viability of independent variate route

During an email conversation with Dave, he suggested that the convolution route might be difficult to implement in terms of processing requirements, and

suggested instead that in the case of the MRC system, the independence of the timing error and fading statistics of each receive channel means that a “joint” PDF describing the probability of all timing and fading variables can be constructed by taking the product of each timing and fading PDF. Therefore we can derive a Gram-Charlier distribution conditioned on a particular set of timing offsets and fading factors, calculate the BER of this theoretical system, and average this over the distribution of timing offsets and fading factors described by the joint PDF.

13/04/14 - MRC adaptation of Gram-Charlier distribution

I went over the process described yesterday in more detail with Dave, and began to implement it by replacing all the $g_k(\Delta)$ terms in his implementation of the non-fading Gram-Charlier series with $\alpha_1^2 g_k(\Delta_1) + \alpha_2^2 g_k(\Delta_2)$ to reflect the effects of fading with a 2 antenna system. In this case, the expected (synchronized) DRB which was previously simply $2g_0(0) = 2$ is now given by $2(\alpha_1^2 g_0(0) + \alpha_2^2 g_0(0)) = 2(\alpha_1^2 + \alpha_2^2)$.

Testing this model, I noticed oscillations at the tails of the distributions which made it impossible to calculate the optimum DRBs, as there were multiple PDF crossings.

14/04/14 - Explaining the oscillations

I raised the issue with Dave, and it turned out that I had forgotten the effects of fading on the channel Gaussian noise. The variance of this noise is scaled by the effects of noise, in this case by $(\alpha_1 + \alpha_2)$. Making this correction I was able to determine the optimum DRB to be equal to $2(\alpha_1^2 g_0(\Delta_1) + \alpha_2^2 g_0(\Delta_2))$, as expected.

Week 13 Summary

In this week I examined Dave’s derivation and modeling of the non-fading case using the Gram-Charlier series, and make a start on my own piece of the analytical work. I successfully extended his derivation to describe the decoder input PDF for a 2 antenna MRC system given known fading statistics and timing offsets, which puts me in a good position to extend this to a more thorough analytical exploration of receiver performance with diversity.

Goals for Week 14

- Evaluate the BER for given fading and timing statistics.
- Average the BER over fading and timing offset distributions to determine the average BER for a given system.

Week 14

16/02/14 - Bit Error Rate Calculation

I implemented the BER calculation for fixed fading and timing statistics as two integrals over the Gram-Charlier PDF approximations, using a given decision region boundary to bound the integrals. Using this I was able to determine the BERs using both the sub-optimum decision region boundaries $2(\alpha_1^2 + \alpha_2^2)$ and the optimum decision region boundaries calculated using last week's code.

18/02/14 - Averaging BER over Tikhonov and Rayleigh distributions

The code was extended to average the BER for each timing error and channel gain over Tikhonov and Rayleigh distributions, respectively, using Gauss-Legendre Quadrature. In this way the average BER for given timing offset and fading statistics can be calculated. Progress was slow as some of the previous code has to be rewritten to fit with its new application. Gauss-Legendre Quadrature was found to be unsuited to the Rayleigh distribution, for reasons I am not yet sure of, so a discrete sampling of the distribution was used instead. At the end of the day the *FindRoot* method used to determine the optimum decision region boundaries refused to converge for certain ranges of timing error and channel gain, so more work must be done to finish it off.

19/02/14 - Root-finding tweaks

I noted that the *FindRoot* method had difficulty converging for both very high and very low channel gains, as the optimum decision region boundary in these cases are much higher or lower than usual. Since the nature of the Rayleigh distribution dictates that these extreme gains are still likely, I used the observation that the gain can be roughly given by $2(\alpha_1^2 + \alpha_2^2)$ for low timing error offsets, and instructed the *FindRoot* method to start the search at this point. This improved both the speed and possibility of converging.

The algorithm was still having difficulty converging for any significant timing error offset, so I greatly reduced the accuracy requirements of the root finding algorithm, under the understanding that we are dealing with average error rates and the goal is to show the rough gain in performance. This brought the algorithm under control, but nonetheless some issues remained:

- For very large timing error offsets, the algorithm converged to an unlikely value. This is unimportant, as in these cases the system has failed to perform any sort of detection.

- The response is somewhat oscillatory. This could be due to the accuracy of the Gram-Charlier series, or less likely the *FindRoot* method. I must ask Dave if he has any idea of how to fix this.
- Even more complexing, the algorithm seems to fail to converge for certain combinations of timing error offset. At first glance it looks like there is a pole or zero at these points. Again, perhaps Dave could suggest a fix for these points.

Insert timing error 2D density plot

20/02/14 - Discussion of Gram-Charlier accuracy

I discussed the issue described above with Dave and Colin, and Dave suggested increasing that it could be a machine precision problem, and that increasing the precision of the variables used in the Gram-Charlier series calculation could solve the problem. Alternatively, Colin suggested that an acceptable answer could be obtained by truncating the sweep to remove the areas where the *FindRoot* algorithm does not converge.

21/02/14 - Overcoming the timing error problems

Through some exploration, the random non-convergent regions were found to disappear at higher AWGN variances. I therefore rose the SNR to 20dB, and found that the oscillations in the root estimates also disappeared.

22/02/14 - Further problems for high timing offsets

I observed that the root finding algorithm failed to converge for large timing error offsets. Noting that it definitely converged for timing offsets of 0.3 and below, and that timing offsets greater than this were extremely unlikely, I decided to truncate the timing offset averaging to $-0.3 \leq \Delta \leq 0.3$.

Plot of roots verses timing error

After implementing this I discovered that the root-finding algorithm results still became increasingly inaccurate for higher timing offsets. This will require further investigation to determine its cause, as left untouched it could significantly reduce the accuracy of the model.

Week 14 Summary

This week saw the implementation of Bit Error Rate calculations for given fading and timing variates, and the averaging of these over the timing and

fading probability distributions was the final step in the implementation of a mathematical model for the described MRC system and allowed a figure for the average BER of the system to be determined. While testing the corners of the implementation, it was found that the analytical approximation failed at certain timing offsets at low SNR values, but this was ignored as larger SNR values are of interest to us. Additionally the optimum DRB finding algorithm does not converge for higher timing error offsets, leading to the decision to ignore the statistically unlikely larger timing offsets. More work will have to be done on reducing the error of the algorithm at middle timing offsets, however.

Goals for Week 15

- Improve the accuracy of the model for large timing error offsets.

Week 15

24/02/14 - Root-finding-defying wiggles

I found that the source of the root-finding errors is in “wiggles” near the roots of the $f_1(y) = f_3(y)$ function used to find the optimum DRB. These seem to be caused by inaccuracies in the Gram-Charlier approximation. Increasing the numerical accuracy did not seem to have any difference, and increasing the length of the Gram-Charlier series only made the oscillations much worse.

25/02/14 - Correct optimum DRB measurements

A brief email exchange with Dave helped me understand the source of the problem. In the model I had built, it was assumed that the instantaneous channel gains and timing offsets were known, and used to make an estimate of the optimum DRB, and the resulting BER was averaged over the fading and timing error PDFs to obtain the average BER. However as mentioned earlier, while the instantaneous channel gains are known in a MRC system (but not an EGC system!), the timing offsets aren't. Hence, the correct approach is to obtain an average PDF for each pair of channel gains, deduce the optimum DRB and BER for each of those channel gains, and then average over the fading and timing error PDFs. Therefore the “wiggles” seen before will be reduced when averaging over the Tikhonov, and shouldn't be present when the optimum DRB is calculated.

I implemented the averaging of the GC PDFs over the Tikhonov timing offset distributions to provide an overall PDF for given channel gains and Tikhonov variances, and from these determined the optimum decision region boundaries. These are now used to determine the BER for different channel gains and timing




Figure 15: “Wiggles” near the roots

offsets, which are averaged over the Rayleigh and Tikhonov distributions to determine the average BER.

26/02/14 - Integration

After a suggestion from Dave, I tried implementing a closed-form solution for the integration of the Gram-Charlier PDF approximation, given solutions for the integration of other identities. In this manner I hoped to significantly reduce the time to determine the average BER.

Starting with the Gram-Charlier definition described earlier,

$$f_{X|\Delta}(y) = \frac{1}{\sigma_X} \phi\left(\frac{y - \mu_X}{\sigma}\right) + \sum_{m=2}^M \frac{\alpha_{2m}}{(2m)! \sigma_X^{2m}} \left[\frac{1}{\sigma_X} \phi\left(\frac{y - \mu_X}{\sigma}\right) H_{2m}\left(\frac{y - \mu_X}{\sigma_X}\right) \right]$$

and incorporating the given identities,

$$\int_{-\infty}^x \frac{1}{\sigma} \phi\left(\frac{y - \mu}{\sigma}\right) dy = \frac{1}{2} \left(1 + \operatorname{erf}\left(\frac{x - \mu}{\sqrt{2}\sigma}\right) \right)$$

$$\int_{-\infty}^x \frac{1}{\sigma} \phi\left(\frac{y-\mu}{\sigma}\right) H_m\left(\frac{y-\mu}{\sigma}\right) dy = -\phi\left(\frac{x-\mu}{\sigma}\right) H_{m-1}\left(\frac{x-\mu}{\sigma}\right)$$

the solution for the integral of the Gram-Charlier series was found to be:

$$\int_{-\infty}^x f_{X|\Delta}(y) dy = \frac{1}{2} \left(1 + \operatorname{erf}\left(\frac{x-\mu_X}{\sqrt{2}\sigma_X}\right) \right) - \sum_{m=2}^M \frac{\alpha_{2m}}{(2m)! \sigma_x^{2m}} \phi\left(\frac{x-\mu_X}{\sigma_X}\right) H_{2m-1}\left(\frac{x-\mu_X}{\sigma_X}\right)$$

I found that while the resulting equation approximated the result found using numerical integration (*NIntegrate*), the two did not match exactly. An interesting question is which correct. My own gut feeling is that the value found using *NIntegrate* is correct, as the algorithm converges to a very precise result and I have very little confidence in my own mathematical abilities, so for the moment I will stick with it. Then again, it is possible that numerical integration is reflecting some unforeseen inaccuracies in my implementation of the Gram-Charlier series, and the closed-form solutions, being mathematically-based, are accurate.

Week 15 Summary

In week 15 I fixed a significant inaccuracy in the model I had built to determine the average BER, stemming from a misconception I had acquired from getting caught up with implementation and forgetting the top-level picture. The numerical implementation of the analytical model seems to be close to finished, and a few tentative tries showed poorer results than those found through simulation. Hopefully work next week will produce the definitive model alongside some quantitative results.

Goals for Week 16

- Complete the implementation of the analytical model.
- Try a few test runs to determine the performance of the standard and modified MRC system.

Week 16

03/03/14 - Rewriting analytical implementation

My machine turned off on Friday afternoon, and I returned from the weekend to find that my *Mathematica* notebook had been corrupted. After some time

trying to manually fix the corrupted file I decided the task was too large and reverted to a week-old backup of the file. As much work had been done in the time between I took the opportunity to rewrite much of the implementation to make it clearer.

04/03/14 - Strengthening the implementation

Some time was spent improving the speed of execution of the implementation and testing it at intermediate points to ensure the results given made intuitive sense. Currently the implementation takes roughly 25mins to run for 5 timing error and channel gain points per channel.

05/03/14 - Assessing the accuracy of the implementation

I spent some more time tweaking the accuracy and precision parameters of the *FindRoot* and *NIntegrate* function to reduce the speed of execution without compromising the accuracy of the system. By looking through the intermediate results of the implementation I found that inaccuracies in the Gram-Charlier approximation for high timing offsets are creating such large error rate estimations for the traditional decision region boundaries that their reduced weighting is insufficient to prevent their effects appearing in the averaged BER value. Unfortunately I cannot imagine how this could be mitigated without ignoring more than the most trivial timing offsets, which wouldn't allow us to show off the full utility of our method.

I added memory to the Gram-Charlier distribution definition to try to speed up the implementation a little more, and started a sweep of timing error variances.

Week 17

14/03/14 - Pre-open day progress

The past weekend was spent trying to account for differences between the analytical and simulation models. Both successfully demonstrated reduced optimum decision region boundaries in the presence of timing errors; however the reduction is much more severe in the simulation than using the analytically derived results. Since I was hoping to be able to derive the optimum DRB values analytically, this was and still remains a major stumbling block.

Dave suggested a modification to the *FindRoot* method which mixes Newton's Method and the Bisection Method in order to prevent divergence, however since this isn't currently an issue I doubt I'll have time to implement it.

Week 18

17/03/14 - Comparing Analytic and Simulation results side-by-side

I spent the weekend combining the conditional analytical and simulation code to provide comparisons between both methods. I found that the analytic solution grossly underestimated the spread of the received symbol, suggesting an issue with the SNR calculations. I was able to confirm that while the analytic method broke down somewhat at higher timing offsets, the response remained somewhat approximate to the pdf generated through simulation. Overall, it did demonstrate similar characteristics to the simulation, just not to the same extent.

Suspecting that this meant that there was a simple mistake in calculating the parameters of the Gram-Charlier approximation, I made a simplified version and asked Dave to look over it.

18/03/14 - Fixing the Gram-Charlier code

Dave got back to me with two mistakes in the code:

- The variable L represents the number of symbols (4 in the case of 4-PAM), not the number of diversity branches as I had mistakenly believed.
- The SNR in the case of channel fading is gained by the RMS sum of the channel gains ($\sqrt{\sum \alpha_i^2}$) and not the sum of the gains ($\sum \alpha_i$).

Implementing these changes, I saw instant improvements in the correlation between the analytical and simulation PDFs. I restarted both the standalone realistic simulation and the conditional simulation/analytical for 2 diversity branches with the changes made to see how closely both methods match.